

Evaluation Report

Criteria

Criterion	Rating	Explanation
Answered problem sections	5/5	This solution addresses the original question of predicting whether a film will receive an IMDB rating of 8 or higher using information that is available before production. You've done a great job excluding any information only available after release like Meta_score, No_of_Votes, and Gross. The limitations mean that the original business question can't be fully answered since the dataset includes IMDB's top 1000 films, but this is well addressed.
Technique judgement	5/5	A supervised classification approach was a great choice for this, since there is a specific target outcome of whether an IMDB rating is greater than 8. There is a good comparison between logistic regression and random forest, as you've explored a more interpretable linear model versus a more flexible model.
Evaluation/interpretation of quality	4/5	You've done a great job with evaluation by looking at precision, recall, F1, ROC-AUC, confusion matrices, and different classification thresholds. One thing I'd look at is choosing a threshold of 0.40 using the test set might weaken the reliability of the reported final performance because the test set is being used for model decisions rather than only the final evaluation. 0.40 was a good choice but the testing and report occurs after the model has seen the test results. I think this would be fine to keep, but if you feel like changing this, it might be cleaner to do choose threshold -> test once -> report performance
Communication	5/5	The notebook is very strong at communicating your decision making to the client and you've explained what the tool does, how it should be used, and its limitations. It's very clear what's being done here!

Open Ended Feedback

This is a well structured and communicated solution to the original problem. It does a good job of connecting the data science work to the original business problem. It also does a great job of recognizing the limitation that the dataset is not representative of the full range of movies. The decision to remove the decade from the model was a good decision, since this feature can't help accurately predict IMDB scores for current and upcoming movies. The only main improvement I'd recommend is to separate model development from final evaluation a bit more strictly. The 0.40 threshold was a good call, but could instead be selected using the training data or crossvalidation then used on the untouched test set after. However, your project shows a great understanding of the overall problem and communicates that well!